

Topics in the Foundations of Artificial Intelligence

A Reader

Draft from May 15, 2024.

Please don't cite or distribute without permission.

Comments welcome!

Levin Hornischer

`Levin.Hornischer@lmu.de`

<https://levinhornischer.github.io/FoundAI/>

Contents

Preface	1
1 Introduction	3
2 Background	5
3 Foundations of symbolic AI	7
4 Standard theory of machine learning	8
5 Refining statistical learning theory	14
6 Computability theory of machine learning	15
7 Using statistical mechanics	16
8 Topological data analysis	17
9 Category theory as a language of machine learning	18
10 Dynamical systems	19
11 Verification of neural networks	20
12 Post-hoc explainability	21
Bibliography	21

Preface

This is the reader for the course *Advanced Topics in the Foundations of AI* given during the summer semester 2024 at *LMU Munich* as part of the *Master in Logic and Philosophy of Science*. The reader is written as the course progresses. A website for the course is found at

<https://levinhornischer.github.io/FoundAI/>.

Comments I'm happy about any comments: spotting typos, finding mistakes, pointing out confusing parts, or simply questions triggered by the material. Just send an informal email to Levin.Hornischer@lmu.de.

Content In recent years, artificial intelligence and, in particular, machine learning made great—but also disconcerting—progress. However, their foundations are, unlike other areas of computer science, less well understood. This situation is sometimes compared to being able to build steam engines without having a theory of thermodynamics.

This seminar is about the mathematical foundation of AI. After a review of the classical theory (Computability Theory, No-Free-Lunch Theorem, Universal Approximation Theorem, etc.), we read some recent research papers to get an overview of some current approaches to the foundations of AI.

Objectives In terms of content, the course aims to convey and overview of the foundations of AI—including both classic material and cutting-edge research. In terms of skills, the course aims to teach the ability to both mathematically and philosophically assess the different approaches to the foundations of AI.

Prerequisites In order to appreciate the literature, the course requires basic familiarity with mathematics (calculus, linear algebra, probability theory), logic (including, ideally, computability theory), and AI (neural networks). Some papers also use more advanced concepts from topology, probability theory, or category theory, so you should also be prepared to read up on those. But they are not assumed: the seminar sessions are,

among others, meant to get clearer on these concepts. Programming skills will of course be useful, but will not be assumed.

Schedule and organization The course is organized as a seminar. Hence, for each session, we have assigned readings, which we then discuss during the session. The reading for each week is announced in the schedule on the course's website. The readings are roughly organized by topic, forming the chapters of this reader.

Background material Some helpful short explainer videos on AI are found [here](#). A excellent mini series on (the mathematics of) neural networks is found [here](#).

Moreover, the material of my companion course on the [Philosophy and Theory of AI](#) might also be helpful. You can take the present course independently of that companion course and vice versa, but they do complement each other. The companion course is more introductory and looks at a broader range of philosophical issues connected to AI and how to theorize about them, while the present course focuses specifically on the more mathematical foundations of neural networks.

A recent edited collection on the mathematical foundations of AI is Grohs and Kutyniok ([2022](#)).

Layout These notes are informal and partially still under construction. For example, there are margin notes to convey more casual comments that you'd rather find in a lecture but usually not in a book. Todo notes indicate, well, that something needs to be done. References are found at the end.

This is a margin note.

This is a todo note

Notation Throughout, 'iff' abbreviates 'if and only if'.

1 Introduction

Summary

We give an outlook of the course and of this document.

The field of AI is typically characterized along the lines of aiming to build “machines that can compute how to act effectively and safely in a wide variety of novel situations” (Russell and Norvig 2021, p. 19). In chapter 2, we briefly collect basic terminology and concepts in AI, to make sure we’re all on the same page.

Regardless of the definition, it is helpful to distinguish two main traditions in AI. They go by varying names, with different connotations depending on the community that uses them, including the following.

1. *Symbolic AI*: classicist, logic-based, Good Old-Fashioned AI (GOFAI), etc.

Example: An algorithm or computer program that, given as input a position in a game of chess, outputs the next best move. This algorithm was written by a programmer.

2. *Subsymbolic AI*: connectionist, non-logicist, machine learning, deep learning, etc.

Example: A neural network that, given as input a pixel image of a handwritten digit, outputs the digit depicted on the image. The neural network was trained to become better at this mapping using thousands of data points, i.e., input images labeled with the digit depicted on them.

(Some might distinguish a third tradition—*statistical AI*—which, in a sense, sits between the two preceding traditions: Like symbolic AI it typically has ‘interpretable’ variables, but they are now continuous random variables, and like subsymbolic AI it typically processes the information in a continuous way.)

For symbolic AI, we have a pretty good theory due to mathematical/philosophical logic and computability theory. We review this in chapter 3. It gives a good idea of what we also would expect of a theory of AI.

For subsymbolic AI—which is most of modern AI—we, however, lack a good theory. So in this course, we mostly focus on approaches to provide such a theory. In chapter 4, we review the main results that the standard theory of machine learning—which is mostly statistical learning theory—can deliver. But we also look at what is still missing for the concrete case of the neural networks that modern AI is built on.

The remaining chapters—as listed in the table of contents—then are about different approaches to fill these gaps in the standard theory of machine learning, or approaches that rethink this theory all together. There is more material than we will be able to cover: especially from the later chapters we will pick the topics based on your interests.

2 Background

Readings

- A textbook introduction to the field of AI: Russell and Norvig (2021, ch. 1).

Key concepts

- History of AI: Ada Lovelace, Alan Turing, McCulloch & Pitts, Logic Theorist, Dartmouth workshop, summers and winters, big data, deep learning revolution.
- Types of AI: symbolic, subsymbolic, statistical
- Definitions of AI: acting humanly (Turing test), thinking humanly (cognitive modeling), thinking rationally (logic, probability), acting rationally (rational agent; perfect vs limited rationality)
- Types of learning tasks: Supervised learning, unsupervised learning, reinforcement learning. Machine learning pipeline (conceptualization, data, model, deployment).
- Key concepts of artificial neural networks: neurons, layers, feed-forward/recurrent, weights, activation function, loss function, backpropagation, learning rate, local/global minima (equilibrium), regularization, overfitting/underfitting.

In this session, we discuss the main reading to get an understanding of each of the key concepts—the basic AI terminology—mentioned above. These key concepts are further illuminated in the additional material mentioned below.

Elaborating on the rational agent idea, Russell and Norvig (2021) write:

In a nutshell, AI has focused on the study and construction of agents that *do the right thing*. What counts as the right thing is defined by the objective that we provide to the agent. This general paradigm is so pervasive that we might call it the *standard*

model. ... [Though,] the standard model assumes that we will supply a fully specified objective to the machine. ... The problem of achieving agreement between our true preferences and the objective we put into the machine is called the *value alignment problem* ... Ultimately, we want agents that are *provably beneficial* to humans (Russell and Norvig 2021, 22–21, emphasis altered).

Further material

- A very accessible overview, written at the beginning of the deep learning revolution: Boden (2016, ch. 1 and 4).
- The Stanford Encyclopedia of Philosophy entry on artificial intelligence: Bringsjord and Govindarajulu (2024).
- Also see the background material mentioned in the preface (explainer videos and companion course).

3 Foundations of symbolic AI

Readings

- A short overview of symbolic AI: Flasiński (2016, ch. 2)
- An overview of computability and complexity theory: Immerman (2021)

Key concepts

- Turing machine
- Church–Turing thesis
- Halting problem
- Entscheidungsproblem
- Tiling problems
- Gödel’s incompleteness theorems
- Computational complexity theory: P vs NP

I might add the contents of the slides that I presented in class

4 Standard theory of machine learning

Readings

- On statistical learning theory and the No-Free-Lunch theorem: Shalev-Shwartz and Ben-David (2014), chapters 2–3 (basic set up), chapter 5 (No-Free-Lunch), sections 6.1–4 (fundamental theorem of statistical learning theory).
- On the approximation theorem: One of the classic papers, Hornik et al. (1989).

Key concepts

- Statistical learning theory
- No-Free-Lunch theorem
- Universal approximation theorem

The statistical learning framework:

- A learner gets input x from some domain X (e.g., a particular papaya) and they need to label this input with a label $h(x) = y$ from the label space Y , here the only labels are 0 (e.g., not tasty) and 1 (e.g., tasty). The learner will see finitely many training data $(x_1, y_1), \dots, (x_m, y_m)$ of input-output pairs, and based on that suggest a general rule $h : X \rightarrow Y$. We now formalize this idea.
- The *domain set* X . Typically the elements are vectors, e.g., $x = (0.1, 0.7)$ saying that the object x is described completely by feature 0 (e.g., the papayas color) having value 0.1 and feature 1 (e.g., the papayas softness) having value 0.7.
- The *label set* Y . Typically $Y = \{0, 1\}$.
- The *training data* $S = \{(x_1, y_1), \dots, (x_m, y_m)\}$ which is a finite subset of $X \times Y$. These are the examples that the learner has access to (e.g., the m -many papayas they have bought, checked the color and softness, and then tasted them to determine whether they were tasty or not).

Elements of S are also called data points or training examples, and S is also called training set.

- The *learner's output*: Based on the training data, the learner has to output a prediction rule, i.e., a rule for how to label points from the domain set. This rule is described as a function $h : X \rightarrow Y$, which is also called predictor, hypothesis, or classifier.
- The *learning algorithm* A : Typically, the learner will follow a general learning algorithm that works not just for the specific training set S but also for other ones. So it is a function A that takes as input finite subsets of $X \times Y$ and outputs a predictor $A(S) : X \rightarrow Y$. One learning algorithm that we introduce below is empirical risk minimization.
- *Sampling*: We assume there is a probability distribution D on the domain set X which describes how likely it is that we see a particular point x (e.g., how likely it is that the learner gets papaya x when going to the market). One says D is a (probabilistic) data-generation model. Importantly, the learner has *no* access to this probability distribution. Rather D describes how the world actually is.
- *True risk*. The error of a classifier h is the probability that it does not predict the correct label. Formally, we assume there is a true labeling function $f : X \rightarrow Y$ (that the learner aims to find) and the error of h is

$$L_{D,f}(h) := D(\{x \in X : h(x) \neq f(x)\}).$$

This is also called generalization error, risk, or true error (to distinguish it from the empirical error/risk that we introduce later). (Below, in the NFL Theorem, we generalize the assumption of a true function f and instead work with a probability distribution over $X \times Y$.)

- *Empirical risk*. The true risk is defined with respect to the distribution D and the true labeling function f , both of which the learner has no access to. The learner can only calculate the error of their predictions on the training dataset:

$$L_S(h) := \frac{1}{m} |\{i \in \{1, \dots, m\} : h(x_i) \neq y_i\}|.$$

- *Empirical Risk Minimization (ERM)* is the learning paradigm of coming up with a predictor h that minimizes the empirical risk $L_S(h)$.

This has to be restricted though: If the learner is allowed to pick any predictor h , they can pick the h which, on input x , predicts y if (x, y) is in the training set and 0 otherwise. This minimizes the empirical risk (it is 0), but it generalizes badly to points x outside of the training set (it all assigns them the same label): one says h *overfits* the data. To avoid this, one fixes a set H of allowed predictors (and the just mentioned h wouldn't usually be allowed). This H is called the *hypothesis class*. The learner has to choose this in advance, before seeing the data. This choice of H is the *inductive bias* (or, positively, prior knowledge) of the learner: they are biased to certain predictors before seeing data (or know a priori that they will better fit the data). (This will lead to the bias-complexity trade-off that we'll discuss later.)

PAC learning:

- What would it mean to say that a learner chose a good hypothesis class H (on domain X and label set Y)? We want that, no matter what the true distribution D and labeling function f are, given a required confidence parameter δ and accuracy parameter ϵ , there is a number of samples $m = m(\delta, \epsilon)$ such that, if we sample m -many examples from D labeled with f , then with confidence $1 - \delta$ the learner knows that they are correct up to ϵ , provided there is a correct hypothesis in the first place. One says: we are *probably approximately correct* (PAC). The formal definition is:
- A hypothesis class H is *PAC-learnable* if there is a function $m_H : (0, 1) \times (0, 1) \rightarrow \mathbb{N}$ and a learning algorithm A such that: For every $\epsilon, \delta \in (0, 1)$, for every probability distribution D over X , and for every labeling function $f : X \rightarrow \{0, 1\}$, if the realizability assumption holds (i.e., there is $h^* \in H$ with $L_{D,f}(h^*) = 0$), then, when running A on $m \geq m(\epsilon, \delta)$ i.i.d. (independently and identically distributed) examples sampled with D and labeled with f , the algorithm return a hypothesis h such that, with probability of at least $1 - \delta$ (over the choice of the examples), $L_{D,f}(h) \leq \epsilon$.
- This can be generalized by dropping the realizability assumption, going beyond the binary label case (multiclass and regression, using more general loss functions).

Though one might ask: why quantify over all D and f and not just those that are likely for the task? E.g., for the papayas some distributions and labelings are more likely than others.

No-Free-Lunch Theorem (NFL Theorem):

- We avoided the overfitting problem of the ERM learning paradigm by making explicit the inductive bias/prior knowledge of the learner in the form of a hypothesis class H . One may ask: is this really necessary, i.e., can't there be a learner who is successful without using any task-specific prior knowledge and can thus solve any task? One way to interpret the NFL theorem is that it answers 'yes': there cannot be such a *universal learner*.
- A bit more precisely, the NFL Theorem takes a learning task to be given by an unknown distribution P over $X \times Y$, and the goal of the learner is to find a predictor $h : X \rightarrow Y$ whose risk $L_P(h)$ is small. As mentioned, this generalizes the framework so far: So far we had a distribution D on X and a true label function f . This determines the distribution P on $X \times Y$ according to which the probability of (x, y) is 0 if $y \neq f(x)$ and otherwise the probability of x according to D . Now we allow any distribution P on $X \times Y$, so we don't assume there is a single true label function, but only a conditional probability of how likely a label is given the input. Accordingly, the loss is

$$L_P(h) = P(\{(x, y) \in X \times Y : h(x) \neq y\},$$

which is also known as *0-1 loss*. The NFL Theorem then says: It is not the case that there is a learning algorithm A and a training set size m such that, for every distribution P over $X \times Y$, if A receives m -many i.i.d. samples from P , there is a high chance it outputs a predictor h that has low risk. In other words, for every learner, there is a task on which it fails, even though another learner succeeds.

- Formally, the NFL theorem is stated as follows. Let A be any learning algorithm for the task of binary classification with respect to the 0-1 loss over a domain X . Let m be any number smaller than $|X|/2$, representing a training set size. Then, there exists a distribution P over $X \times \{0, 1\}$ such that:
 1. There exists a function $f : X \rightarrow \{0, 1\}$ with $L_P(f) = 0$, but
 2. With probability of at least $1/7$ over the choice of $S \sim P^m$ we have that $L_P(A(S)) \geq 1/8$.

So A fails on this task while the ERM learner with hypothesis class $H = \{f\}$ succeeds.

For a discussion of interpretations of the NFL theorem, see the further reading below (Sterkenburg and Grünwald 2021).

Bias-complexity tradeoff

- Since we don't have a universal learner—who performs best possible on any task—, we cannot get around analyzing a learner on a specific task. So we want to get some guarantee how bad the true error of the learner is. As we'll describe now, we can split this error up into two parts, but there is a tradeoff in that improving one part tends to worsen the other, and vice versa.
- *Error decomposition.* We can analyze the error $L_P(A(S))$ of our learning algorithm/ERM predictor A on a given dataset S as $\epsilon_{\text{app}} + \epsilon_{\text{est}}$ with

$$\epsilon_{\text{app}} := \min_{h \in H} L_P(h) \qquad \epsilon_{\text{est}} := L_P(A(S)) - \epsilon_{\text{app}},$$

where ϵ_{app} is the *approximation error* (the minimal risk achievable by a predictor from the hypothesis class) and ϵ_{est} is the *estimation error* (the difference between the empirical minimal risk $L_P(A(S))$ and the true minimal risk ϵ_{app}).

- *Bias-complexity tradeoff.* To reduce the error $L_P(A(S))$ we hence want to reduce both the approximation error and the estimation error. However, this comes at a tradeoff:
 1. To reduce the approximation error, we want a large, more *complex* hypothesis class; but, as we saw, this may lead to *overfitting* (empirical risk is low, but true risk is high), so the estimation error can be high.
 2. To reduce the estimation error, we might rather choose a small, more *biased* hypothesis class, but this results in a high approximation error, because the predictors are now *underfitting* the data.

So learning theory studies how to find rich hypothesis classes H for which we still have reasonable estimation errors. We will further discuss this for neural networks in chapter 5 (the universal approximation theorem will show that—as the name suggests—their approximation error is zero).

VC-dimension and the fundamental theorem of PAC learning:

- How can we ensure our learner has a low error? Reducing the approximation error is a matter of picking the right prior knowledge encapsulated as a hypothesis class. For the estimation error, the key is to realize that PAC learnability bounds the estimation error. So we

Under the realizability assumption, the approximation error is zero, so the estimation error can be bounded by ϵ with probability $1 - \delta$ when sampling a dataset of size $\geq m(\delta, \epsilon)$.

would like to have a guarantee for PAC learnability. This is provided by the so-called VC-dimension, which is defined as follows.

- First, some helpful terminology: If H is a hypothesis class and $C \subseteq X$, then H *shatters* C if $\{h \upharpoonright C : h \in H\} = Y^C$, where $h \upharpoonright C$ is the restriction of the function $h : X \rightarrow Y$ to the set C , and Y^C is the set of all functions from C to Y .
- The *VC-dimension* of a hypothesis class H is the maximal size of a finite set $C \subseteq X$ that can be shattered by H . If H shatters sets of arbitrarily large size, it has infinite VC-dimension.
- The *fundamental theorem*: Let H be a hypothesis class of functions from X to $\{0, 1\}$, and let the loss function be the 0-1 loss. Then the following are equivalent
 1. H is PAC learnable
 2. Any ERM rule is a successful PAC learner for H
 3. H has a finite VC-dimension.

The idea is: If H shatters a set C of size $2m$, then we cannot learn H using m examples: all ways of labeling the other m instances in C are possible according to some hypothesis in H , i.e., prior knowledge doesn't exclude any of those.

The full version of the theorem has some further equivalent conditions.

In fact, if the VC-dimension of H is $d < \infty$, there are constants C_1 and C_2 such that H is PAC learnable with sample complexity

$$C_1 \frac{d + \log(1/\delta)}{\epsilon} \leq m(\epsilon, \delta) \leq C_2 \frac{d \log(1/\epsilon) + \log(1/\delta)}{\epsilon}.$$

Further material

- A modern, more general approach to the universal approximation theorem: Kratsios (2021).
- A discussion of the no-free-lunch theorem: Sterkenburg and Grünwald (2021).

5 Refining statistical learning theory

Readings

- Berner et al. (2022)
- Belkin (2021)

Key concepts

- TBA

Further material

- TBA

6 Computability theory of machine learning

Readings

- Colbrook et al. (2022)
- Caro (2023)

Key concepts

- TBA

Further material

- An older overview: Šíma and Orponen (2003).
- Delétang et al. (2023)

7 Using statistical mechanics

Readings

- Roberts and Yaida (2022)
- Bahri et al. (2020)

Key concepts

- TBA

Further material

- The book by Roberts and Yaida (2022) has been presented in a course (<https://deeplearningtheory.com/lectures/>).

8 Topological data analysis

Readings

- Naitzat et al. (2020)
- Overview: Hensel et al. (2021)

Key concepts

- TBA

Further material

- For a short explanation of persistent homology, see [here](#).
- For a popular science application to neuroscience, see [here](#).

9 Category theory as a language of machine learning

Readings

- Shiebler et al. (2021)
- Bradley et al. (2021)
- Bronstein et al. (2021)
- van Bekkum et al. (2021)

Key concepts

- TBA

Further material

- More on categories for AI in this course (<https://cats.for.ai/>) and this list of papers on the topic (https://github.com/bgavran/Category_Theory_Machine_Learning).

10 Dynamical systems

Readings

- Overview: Bournez and Pouly ([2021](#))

Key concepts

- TBA

Further material

- TBA

11 Verification of neural networks

Readings

- Albarghouthi (2021)

Key concepts

- TBA

Further material

- TBA

12 Post-hoc explainability

Readings

- Bilodeau et al. (2024)

Key concepts

- TBA

Further material

- TBA

Bibliography

- Albarghouthi, A. (2021). *Introduction to Neural Network Verification*. arXiv: [2109.10317 \[cs.LG\]](#) (cit. on p. 20).
- Bahri, Y. et al. (2020). “Statistical Mechanics of Deep Learning.” In: *Annual review of condensed matter physics* 11.1, pp. 501–528. DOI: [10.1146/annurev-conmatphys-031119-050745](#) (cit. on p. 16).
- Belkin, M. (2021). “Fit without fear: remarkable mathematical phenomena of deep learning through the prism of interpolation.” In: *Acta Numerica* 30, pp. 203–248 (cit. on p. 14).
- Berner, J. et al. (2022). “The Modern Mathematics of Deep Learning.” In: *Mathematical Aspects of Deep Learning*. Ed. by P. Grohs and G. Kutyniok. Cambridge: Cambridge University Press, pp. 1–111. DOI: [10.1017/9781009025096.002](#) (cit. on p. 14).
- Bilodeau, B. et al. (2024). “Impossibility theorems for feature attribution.” In: *Proceedings of the National Academy of Sciences* 121.2, e2304406120 (cit. on p. 21).
- Boden, M. A. (2016). *AI: Its nature and future*. Oxford: Oxford University Press (cit. on p. 6).
- Bournez, O. and A. Pouly (2021). “A Survey on Analog Models of Computation.” In: *Handbook of Computability and Complexity in Analysis*. Ed. by V. Brattka and P. Hertling. Cham: Springer International Publishing, pp. 173–226. DOI: [10.1007/978-3-030-59234-9_6](#) (cit. on p. 19).
- Bradley, T.-D., J. Terilla, and Y. Vlassopoulos (2021). *An enriched category theory of language: from syntax to semantics*. arXiv: [2106.07890 \[math.CT\]](#) (cit. on p. 18).
- Bringsjord, S. and N. S. Govindarajulu (2024). “Artificial Intelligence.” In: *The Stanford Encyclopedia of Philosophy*. Ed. by E. N. Zalta and U. Nodelman. Summer 2024. Metaphysics Research Lab, Stanford University (cit. on p. 6).

- Bronstein, M. M. et al. (2021). *Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges*. arXiv: 2104.13478 [cs.LG] (cit. on p. 18).
- Caro, M. C. (12/2023). “From undecidability of non-triviality and finiteness to undecidability of learnability.” In: *International Journal of Approximate Reasoning* 163, p. 109057. DOI: 10.1016/j.ijar.2023.109057. URL: <https://arxiv.org/abs/2106.01382> (cit. on p. 15).
- Colbrook, M. J., V. Antun, and A. C. Hansen (2022). “The difficulty of computing stable and accurate neural networks: On the barriers of deep learning and Smale’s 18th problem.” In: *Proceedings of the National Academy of Sciences* 119.12, e2107151119. DOI: 10.1073/pnas.2107151119. eprint: <https://www.pnas.org/doi/pdf/10.1073/pnas.2107151119>. URL: <https://www.pnas.org/doi/abs/10.1073/pnas.2107151119> (cit. on p. 15).
- Delétang, G. et al. (2023). *Neural Networks and the Chomsky Hierarchy*. arXiv: 2207.02098 [cs.LG] (cit. on p. 15).
- Flasiński, M. (2016). *Introduction to Artificial Intelligence*. Cham: Springer. DOI: <https://doi.org/10.1007/978-3-319-40022-8> (cit. on p. 7).
- Grohs, P. and G. Kutyniok, eds. (2022). *Mathematical Aspects of Deep Learning*. Cambridge University Press. DOI: 10.1017/9781009025096 (cit. on p. 2).
- Hensel, F., M. Moor, and B. Rieck (2021). “A Survey of Topological Machine Learning Methods.” In: *Frontiers in Artificial Intelligence* 4. DOI: 10.3389/frai.2021.681108 (cit. on p. 17).
- Hornik, K., M. Stinchcombe, and H. White (1989). “Multilayer feedforward networks are universal approximators.” In: *Neural Networks* 2.5, pp. 359–366. DOI: [https://doi.org/10.1016/0893-6080\(89\)90020-8](https://doi.org/10.1016/0893-6080(89)90020-8) (cit. on p. 8).
- Immerman, N. (2021). “Computability and Complexity.” In: *The Stanford Encyclopedia of Philosophy*. Ed. by E. N. Zalta. Winter 2021. Metaphysics Research Lab, Stanford University (cit. on p. 7).
- Kratsios, A. (2021). “The Universal Approximation Property.” In: *Annals of Mathematics and Artificial Intelligence* 89.435-469. DOI: <https://doi.org/10.1007/s10472-020-09723-1> (cit. on p. 13).

- Naitzat, G., A. Zhitnikov, and L.-H. Lim (2020). “Topology of Deep Neural Networks.” In: *Journal of Machine Learning Research* 21.184, pp. 1–40. URL: <http://jmlr.org/papers/v21/20-345.html> (cit. on p. 17).
- Roberts, D. A. and S. Yaida (2022). *The Principles of Deep Learning Theory: An Effective Theory Approach to Understanding Neural Networks*. Cambridge: Cambridge University Press. DOI: [10.1017/9781009023405](https://doi.org/10.1017/9781009023405). URL: <https://arxiv.org/abs/2106.10165> (cit. on p. 16).
- Russell, S. J. and P. Norvig (2021). *Artificial Intelligence: A Modern Approach*. Pearson series in artificial intelligence. Harlow: Pearson (cit. on pp. 3, 5 sq.).
- Shalev-Shwartz, S. and S. Ben-David (2014). *Understanding Machine Learning: From Theory to Algorithms*. A copy for personal use only at <https://www.cs.huji.ac.il/~shais/UnderstandingMachineLearning/copy.html>. New York: Cambridge University Press (cit. on p. 8).
- Shiebler, D., B. Gavranović, and P. Wilson (2021). *Category Theory in Machine Learning*. URL: <https://arxiv.org/abs/2106.07032> (cit. on p. 18).
- Šíma, J. and P. Orponen (2003). “General-Purpose Computation with Neural Networks: A Survey of Complexity Theoretic Results.” In: *Neural Computation* 15, pp. 2727–2778 (cit. on p. 15).
- Sterkenburg, T. F. and P. D. Grünwald (2021). “The no-free-lunch theorems of supervised learning.” In: *Synthese* 199.3-4, pp. 9979–10015. DOI: [10.1007/s11229-021-03233-1](https://doi.org/10.1007/s11229-021-03233-1) (cit. on pp. 11, 13).
- Van Bekkum, M. et al. (2021). “Modular design patterns for hybrid learning and reasoning systems: a taxonomy, patterns and use cases.” In: *Applied Intelligence* 51, pp. 6528–6546. DOI: [10.1007/s10489-021-02394-3](https://doi.org/10.1007/s10489-021-02394-3) (cit. on p. 18).